

```

1 !git clone https://github.com/Franck-Dernoncourt/pubmed-rct

Cloning into 'pubmed-rct'...
remote: Enumerating objects: 39, done.
remote: Counting objects: 100% (14/14), done.
remote: Compressing objects: 100% (9/9), done.
remote: Total 39 (delta 8), reused 5 (delta 5), pack-reused 25
Receiving objects: 100% (39/39), 177.08 MiB | 35.46 MiB/s, done.
Resolving deltas: 100% (15/15), done.

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 import tensorflow as tf
7 import tensorflow_hub as hub
8 from tensorflow.keras import layers
9 from tensorflow.keras.utils import plot_model
10
11 from sklearn.preprocessing import OneHotEncoder, LabelEncoder
12 from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

1 def file_data(filepath):
2     with open(filepath) as f:
3         return f.readlines()

1 dataset_directory = "pubmed-rct/PubMed_20k_RCT/"

1 import os
2 os.getcwd()

'/content'

1 # reading data from directory and storing in another directory
2 training_dir = file_data(dataset_directory + 'train.txt')

1 training_dir[13].isspace()

True

1 learning_sample = training_dir[:100]

1 learning_sample[:20]

['###24293578\n',
 'OBJECTIVE\tTo investigate the efficacy of 6 weeks of daily low-dose oral prednisolone in improving pain , mobility , and systemic
 low-grade inflammation in the short term and whether the effect would be sustained at 12 weeks in older adults with moderate to
 severe knee osteoarthritis ( OA ) .\n',
 'METHODS\tA total of 125 patients with primary knee OA were randomized 1:1 ; 63 received 7.5 mg/day of prednisolone and 62
 received placebo for 6 weeks .\n',
 'METHODS\tOutcome measures included pain reduction and improvement in function scores and systemic inflammation markers .\n',
 'METHODS\tPain was assessed using the visual analog pain scale ( 0-100 mm ) .\n',
 'METHODS\tSecondary outcome measures included the Western Ontario and McMaster Universities Osteoarthritis Index scores , patient
 global assessment ( PGA ) of the severity of knee OA , and 6-min walk distance ( 6MWD ) .\n',
 'METHODS\tSerum levels of interleukin 1 ( IL-1 ) , IL-6 , tumor necrosis factor ( TNF ) - , and high-sensitivity C-reactive
 protein ( hsCRP ) were measured .\n',
 'RESULTS\tThere was a clinically relevant reduction in the intervention group compared to the placebo group for knee pain ,
 physical function , PGA , and 6MWD at 6 weeks .\n',
 'RESULTS\tThe mean difference between treatment arms ( 95 % CI ) was 10.9 ( 4.8-18 .0 ) , p < 0.001 ; 9.5 ( 3.7-15 .4 ) , p < 0.05
 ; 15.7 ( 5.3-26 .1 ) , p < 0.001 ; and 86.9 ( 29.8-144 .1 ) , p < 0.05 , respectively .\n',
 'RESULTS\tFurther , there was a clinically relevant reduction in the serum levels of IL-1 , IL-6 , TNF - , and hsCRP at 6 weeks in
 the intervention group when compared to the placebo group .\n',
 'RESULTS\tThese differences remained significant at 12 weeks .\n',
 'RESULTS\tThe Outcome Measures in Rheumatology Clinical Trials-Osteoarthritis Research Society International responder rate was 65
 % in the intervention group and 34 % in the placebo group ( p < 0.05 ) .\n',
 'CONCLUSIONS\tLow-dose oral prednisolone had both a short-term and a longer sustained effect resulting in less knee pain , better
 physical function , and attenuation of systemic inflammation in older patients with knee OA ( ClinicalTrials.gov identifier
 NCT01619163 ) .\n',
 '\n',
 '###24854809\n',
 'BACKGROUND\tEmotional eating is associated with overeating and the development of obesity .\n',
 'BACKGROUND\tYet , empirical evidence for individual ( trait ) differences in emotional eating and cognitive mechanisms that
 contribute to eating during sad mood remain equivocal .\n',
 'OBJECTIVE\tThe aim of this study was to test if attention bias for food moderates the effect of self-reported emotional eating
 during sad mood ( vs neutral mood ) on actual food intake .\n',
 'OBJECTIVE\tIt was expected that emotional eating is predictive of elevated attention for food and higher food intake after an
 experimentally induced sad mood and that attentional maintenance on food predicts food intake during a sad versus a neutral mood
 .\n',

```

```
'METHODS\tParticipants ( N = 85 ) were randomly assigned to one of the two experimental mood induction conditions ( sad/neutral )
.\n']
```

```
1 def get_data(path):
2     raw_data = file_data(path)
3     article_data = ''
4     data_sample = []
5     article_id = 0
6
7     for line in raw_data:
8         if line.startswith('###'):
9             article_id = int(line.replace('###', '').replace('\n', '')) # getting the ID of the blog
10            article_data = ''
11
12            elif line.isspace(): # end of one article
13                article_data_split = article_data.splitlines()
14
15                for article_line_number, article_line in enumerate(article_data_split):
16                    line_data = {} # storing in dictionary
17                    target_text_split = article_line.split("\t") # split target label from text
18
19                    line_data["article_id"] = article_id # storing article ID
20                    line_data["line_id"] = f'{article_id}_{article_line_number}_{len(article_data_split)}' # create id for each line from
21                    line_data["abstract_text"] = target_text_split[1] # storing target text
22                    line_data["line_number"] = article_line_number # get line number in the article
23                    line_data["total_lines"] = len(article_data_split) # total number of lines in article
24                    line_data['current_line'] = f'{article_line_number}_{len(article_data_split)}' # embedding article length with line
25                    line_data["target"] = target_text_split[0] # storing the label of the target
26
27                    data_sample.append(line_data) # appending the sample data
28
29            else: #line into article line if the condition is not satisfied.
30                article_data += line
31
32    return data_sample

```

```
1 train_samples = get_data(dataset_directory + "train.txt")
2 val_samples = get_data(dataset_directory + "dev.txt")
3 test_samples = get_data(dataset_directory + "test.txt")

```

```
1 train_samples
```



```

'abstract_text': 'The aim of this study was to test if attention bias for food moderates the effect of self-reported
emotional eating during sad mood ( vs neutral mood ) on actual food intake .',
'line_number': 2,
'total_lines': 11,
'current_line': '2_11',
'target': 'OBJECTIVE'},
{'article_id': 24854809,
'line_id': '24854809_3_11',
'abstract_text': 'It was expected that emotional eating is predictive of elevated attention for food and higher food intake
after an experimentally induced sad mood and that attentional maintenance on food predicts food intake during a sad versus a
neutral mood .',
'line_number': 3,
'total_lines': 11,
'current_line': '3_11',
'target': 'OBJECTIVE'},
{'article_id': 24854809,
'line_id': '24854809_3_11',
'abstract_text': 'It was expected that emotional eating is predictive of elevated attention for food and higher food intake
after an experimentally induced sad mood and that attentional maintenance on food predicts food intake during a sad versus a
neutral mood .',
'line_number': 3,
'total_lines': 11,
'current_line': '3_11',
'target': 'OBJECTIVE'}

1 # length of the samples
2 print('1. lenth of the train samples after processing: ',len(train_samples))
3 print('2. lenth of the test samples after processing: ',len(test_samples))
4 print('3. lenth of the val samples after processing: ',len(val_samples))

1. lenth of the train samples after processing: 180040
2. lenth of the test samples after processing: 30135
3. lenth of the val samples after processing: 30212

1 train_dataframe = pd.DataFrame(train_samples)
2 val_dataframe = pd.DataFrame(val_samples)
3 test_dataframe = pd.DataFrame(test_samples)

1 train_dataframe['target'].value_counts().sort_values()

OBJECTIVE      13839
BACKGROUND     21727
CONCLUSIONS  27168
RESULTS        57953
METHODS        59353
Name: target, dtype: int64

1 train_dataframe['target'].value_counts().values


array([59353, 57953, 27168, 21727, 13839])

1 train_dist = train_dataframe['target'].value_counts().sort_values()
2 val_dist = val_dataframe['target'].value_counts().sort_values()
3 test_dis = test_dataframe['target'].value_counts().sort_values()

1 print(plt.style.available)

['Solarize_Light2', '_classic_test_patch', '_mpl-gallery', '_mpl-gallery-nogrid', 'bmh', 'classic', 'dark_background', 'fast', 'five
fig, ax = plt.subplots(3, 1, figsize=(12, 12))
2 #plt.style.use('grayscale')
3 sns.barplot(x=train_dist.values, y=list(train_dist.index), ax=ax[0],orient='h')
4 ax[0].set_title('Train distribution over labels')
5
6 sns.barplot(x=val_dist.values, y=list(val_dist.index), ax=ax[1],orient='h')
7 ax[1].set_title('validation distribution over labels')
8
9 sns.barplot(x=test_dis.values, y=list(test_dis.index), ax=ax[2],orient='h')
10 ax[1].set_title('Test distribution over labels')
11 plt.subplots_adjust(hspace=0.5)
12
13 plt.show()
14

```



```
1 ohc = OneHotEncoder(sparse=False)
```

```
1 train_labels_one_hot_encoded = ohc.fit_transform(train_dataframe["target"].to_numpy().reshape(-1, 1))
```

```
2 val_labels_one_hot_encoded = ohc.transform(val_dataframe["target"].to_numpy().reshape(-1, 1))
```

```
3 test_labels_one_hot_encoded = ohc.transform(test_dataframe["target"].to_numpy().reshape(-1, 1))
```

/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_encoders.py:868: FutureWarning: `sparse` was renamed to `sparse_output` in 1.0; the `sparse_output` parameter will be removed in 1.2; the new default behavior is to output dense arrays, to retain the old behavior use `sparse_output=True`
warnings.warn(

```
1 LE = LabelEncoder()
```

```
1 train_labels_encoded = LE.fit_transform(train_dataframe["target"].to_numpy())
```

```
1 val_labels_encoded = LE.transform(val_dataframe["target"].to_numpy())
```

```
2 test_labels_encoded = LE.transform(test_dataframe["target"].to_numpy())
```

```
1 train_labels_encoded
```

array([3, 2, 2, ..., 4, 1, 1])

```

1 # number of classes
2 print('Number of classes: ', len(LE.classes_))
3 print('All the Classes :', LE.classes_)

→ Number of classes: 5
   All the Classes : ['BACKGROUND' 'CONCLUSIONS' 'METHODS' 'OBJECTIVE' 'RESULTS']

1 ohc_for_lines = OneHotEncoder(sparse=False)

1 train_lines_encoded = ohc_for_lines.fit_transform(train_dataframe["current_line"].to_numpy().reshape(-1, 1)).astype(np.float32)
2 val_lines_encoded = ohc_for_lines.transform(val_dataframe["current_line"].to_numpy().reshape(-1, 1)).astype(np.float32)
3 test_lines_encoded = ohc_for_lines.transform(test_dataframe["current_line"].to_numpy().reshape(-1, 1)).astype(np.float32)

→ /usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_encoders.py:868: FutureWarning: `sparse` was renamed to `sparse_output`
   warnings.warn(

```



```

1 def apply_smoothing(line_one_hot, esp=0.01):
2     return np.abs(line_one_hot - esp)
3
4 def revert_smoothing(line_one_hot_smooth, esp=0.01):
5     unsmooth = line_one_hot_smooth
6
7     unsmooth[unsmooth>esp] = 1.0
8     unsmooth[unsmooth<=esp] = 0.0
9
10    return unsmooth

1 # applying smoothing on train lines
2 train_lines_encoded = apply_smoothing(train_lines_encoded)

1 def create_pipeline(features, labels, batch_size=32, shuffle=False, cache=False, prefetch=False) -> tf.data.Dataset:
2
3     AUTOTUNE = tf.data.AUTOTUNE
4     ds = tf.data.Dataset.from_tensor_slices((features, labels))
5
6     # shuffling on the basis of condition
7     if shuffle:
8         ds = ds.shuffle(buffer_size=1000)
9         #batches
10    ds = ds.batch(batch_size)
11
12    # Apply caching based on condition
13    if cache:
14        ds = ds.cache(buffer_size=AUTOTUNE)
15    if prefetch:
16        ds = ds.prefetch(buffer_size=AUTOTUNE)
17    return ds

1 train_sentence_list = train_dataframe["abstract_text"].tolist()
2 val_sentence_list = val_dataframe["abstract_text"].tolist()
3 test_sentence_list = test_dataframe["abstract_text"].tolist()

1 BATCH_SIZE = 32
2
3 # Create preprocessed training dataset
4 train_features = (train_sentence_list, train_lines_encoded.astype(np.float32))
5 train_labels = train_labels_one_hot_encoded.astype(np.float32)
6
7 # creating pipeline
8 train_ds = create_pipeline(
9     train_features, train_labels,
10    batch_size=BATCH_SIZE, shuffle=True,
11    cache=False, prefetch=True)

1 # applying the same to val_data
2 val_features = (val_sentence_list, val_lines_encoded.astype(np.float32))
3 val_labels = val_labels_one_hot_encoded.astype(np.float32)
4
5 val_ds = create_pipeline(val_features, val_labels,
6     batch_size=BATCH_SIZE, shuffle=False,
7     cache=False, prefetch=True)

```

```

1 # applying to test data
2 test_features = (test_sentence_list, test_lines_encoded.astype(np.float32))
3 test_labels = test_labels_one_hot_encoded.astype(np.float32)
4
5 test_ds = create_pipeline(test_features, test_labels,
6                           batch_size=BATCH_SIZE, shuffle=False,
7                           cache=False, prefetch=True)

1 # eetting model/processor from hub
2 def get_tfhub_model(model_link, model_name, model_trainable=False):
3     return hub.KerasLayer(model_link,
4                             trainable=model_trainable,
5                             name=model_name)

1 encoder_link = 'https://tfhub.dev/google/universal-sentence-encoder/4'
2 encoder_name = 'universal_sentence_encoder'
3 encoder_trainable=False # set trainable to False for inference-only
4
5 encoder = get_tfhub_model(encoder_link, encoder_name, model_trainable=encoder_trainable)

1 class SelfAttentionBlock(layers.Layer):
2
3     def __init__(self, units: int, activation='gelu', kernel_initializer='GlorotNormal', **kwargs):
4         super(SelfAttentionBlock, self).__init__(**kwargs)
5         self.units = units
6         self.activation = activation
7         self.kernel_initializer = tf.keras.initializers.deserialize(kernel_initializer)
8
9         self.query = layers.LSTM(self.units, activation=self.activation,
10                                  kernel_initializer=self.kernel_initializer,
11                                  return_sequences=True, name=f'block_query_lstm')
12
13         self.value = layers.LSTM(self.units, activation=self.activation,
14                                   kernel_initializer=self.kernel_initializer, go_backwards=True,
15                                   return_sequences=True, name=f'block_value_lstm')
16
17         self.attention = layers.AdditiveAttention(name='block_attention')
18         self.average_pooler = layers.GlobalAveragePooling1D(name='block_average_pooler')
19         self.query_batch_norm = layers.BatchNormalization(name='block_query_batch_norm')
20         self.attention_batch_norm = layers.BatchNormalization(name='block_attention_batch_norm')
21         self.residual = layers.Add(name='block_residual')
22
23
24     def __call__(self, x):
25         dim_expand_layer = layers.Lambda(lambda embedding: tf.expand_dims(embedding, axis=1), name='block_dim_expand')
26         x_expanded = dim_expand_layer(x)
27
28         #LSTM sequences
29         block_query = self.query(x_expanded)
30         block_value = self.value(x_expanded)
31
32         #self-attention to LSTM
33         block_attention = self.attention([block_query, block_value])
34
35         # Apply GlobalAvgPooling and BatchNorm to ensure output shape is 1D
36         block_query_pooling = self.average_pooler(block_query)
37         block_query_batch_norm = self.query_batch_norm(block_query_pooling)
38
39         block_attention_pooling = self.average_pooler(block_attention)
40         block_attention_batch_norm = self.attention_batch_norm(block_attention_pooling)
41
42         # Generate addition residual with processed query and attention
43         block_residual = self.residual([block_query_batch_norm, block_attention_batch_norm])
44
45         return block_residual

```

```

1 def build_model():
2     abstract_input = layers.Input(shape=[], dtype=tf.string, name='abstract_text_input')
3     abstract_current_line = layers.Input(shape=(460), dtype=tf.float32, name='abstract_current_line')
4
5     initializer = tf.keras.initializers.GlorotNormal()
6
7     abstract_embedding = encoder(abstract_input)
8     add_attention_block = SelfAttentionBlock(64)(abstract_embedding)
9     abstract_dense = layers.Dense(64, activation='gelu', kernel_initializer=initializer)(abstract_embedding)
10    attention_residual = layers.Multiply(name='mul_residual')([add_attention_block, abstract_dense])
11    current_line_dense = layers.Dense(32, activation='gelu', kernel_initializer=initializer)(abstract_current_line)
12    current_line_dense = layers.Dropout(0.2)(current_line_dense)
13
14
15    streams_concat = layers.Concatenate()([
16        attention_residual,
17        current_line_dense
18    ])
19
20    output_layer = layers.Dense(64, activation='gelu', kernel_initializer=initializer)(streams_concat)
21    output_layer = layers.Dense(5, activation='softmax', kernel_initializer=initializer)(output_layer)
22
23    return tf.keras.Model(inputs=[abstract_input,
24                                abstract_current_line],
25                           outputs=[output_layer], name="use_attention_model")

```

```
1 model = build_model()
```

⚡ /usr/local/lib/python3.10/dist-packages/keras/src/initializers/initializers.py:120: UserWarning: The initializer GlorotNormal is uns

warnings.warn(

```
1 model.summary()
```

⚡ Model: "use_attention_model"

Layer (type)	Output Shape	Param #	Connected to
abstract_text_input (Input Layer)	[(None,)]	0	[]
universal_sentence_encoder (KerasLayer)	(None, 512)	2567978 24	['abstract_text_input[0][0]']
block_dim_expand (Lambda)	(None, 1, 512)	0	['universal_sentence_encoder[0][0]']
block_query_lstm (LSTM)	(None, 1, 64)	147712	['block_dim_expand[0][0]']
block_value_lstm (LSTM)	(None, 1, 64)	147712	['block_dim_expand[0][0]']
block_attention (AdditiveAttention)	(None, 1, 64)	64	['block_query_lstm[0][0]', 'block_value_lstm[0][0]']
block_average_pooler (GlobalAveragePooling1D)	(None, 64)	0	['block_query_lstm[0][0]', 'block_attention[0][0]']
block_query_batch_norm (BatchNormalization)	(None, 64)	256	['block_average_pooler[0][0]']
block_attention_batch_norm (BatchNormalization)	(None, 64)	256	['block_average_pooler[1][0]']
abstract_current_line (Input Layer)	[(None, 460)]	0	[]
block_residual (Add)	(None, 64)	0	['block_query_batch_norm[0][0]', 'block_attention_batch_norm[0][0]']
dense (Dense)	(None, 64)	32832	['universal_sentence_encoder[0][0]']
dense_1 (Dense)	(None, 32)	14752	['abstract_current_line[0][0]']
mul_residual (Multiply)	(None, 64)	0	['block_residual[0][0]', 'dense[0][0]']
dropout (Dropout)	(None, 32)	0	['dense_1[0][0]']
concatenate (Concatenate)	(None, 96)	0	['mul_residual[0][0]', 'dropout[0][0]']

```

dense_2 (Dense)          (None, 64)          6208      ['concatenate[0][0]']

dense_3 (Dense)          (None, 5)           325       ['dense_2[0][0]']

```

```

=====
Total params: 257147941 (980.94 MB)
Trainable params: 240861 (1.22 MB)

```

```

1 def train_model(model, num_epochs, callbacks_list, tf_train_data,
2                 tf_valid_data=None, shuffling=False):
3
4     model_history = {}
5
6     if tf_valid_data != None:
7         model_history = model.fit(tf_train_data,
8                                   epochs=num_epochs,
9                                   validation_data=tf_valid_data,
10                                  validation_steps=int(len(tf_valid_data)),
11                                  callbacks=callbacks_list,
12                                  shuffle=shuffling)
13
14     if tf_valid_data == None:
15         model_history = model.fit(tf_train_data,
16                                   epochs=num_epochs,
17                                   callbacks=callbacks_list,
18                                   shuffle=shuffling)
19
20     return model_history

```

```

1 early_stopping_callback = tf.keras.callbacks.EarlyStopping(
2     monitor='val_loss',
3     patience=4,
4     restore_best_weights=True)
5
6 reduce_lr_callback = tf.keras.callbacks.ReduceLROnPlateau(
7     monitor='val_loss',
8     patience=2,
9     factor=0.1,
10    verbose=1)
11
12 EPOCHS = 10
13 CALLBACKS = [early_stopping_callback, reduce_lr_callback]

```

```

1 train_sentences_count = len(train_sentence_list)
2 val_sentences_count = len(val_sentence_list)
3 test_sentences_count = len(test_sentence_list)
4 total_sentences_count = train_sentences_count + val_sentences_count + test_sentences_count

```

```

1 tf.random.set_seed(42)
2
3 model.compile(
4     loss=tf.keras.losses.CategoricalCrossentropy(label_smoothing=0.1),
5     optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
6     metrics=['accuracy'])
7 )
8
9 print(f'Training {model.name}.')
10 print(f'Train on {train_sentences_count} samples, validate on {val_sentences_count} samples.')
11 print('-----')
12
13 model_history = train_model(
14     model, EPOCHS, CALLBACKS,
15     train_ds, val_ds,
16     shuffling=False)
17 )

```

```

🔄 Training use_attention_model.
Train on 180040 samples, validate on 30212 samples.
-----
Epoch 1/10
5627/5627 [=====] - 145s 24ms/step - loss: 0.6963 - accuracy: 0.8476 - val_loss: 0.6421 - val_accuracy: 0.8
Epoch 2/10
5627/5627 [=====] - 129s 23ms/step - loss: 0.6315 - accuracy: 0.8794 - val_loss: 0.6390 - val_accuracy: 0.8
Epoch 3/10
5627/5627 [=====] - 142s 25ms/step - loss: 0.6140 - accuracy: 0.8878 - val_loss: 0.6339 - val_accuracy: 0.8
Epoch 4/10
5627/5627 [=====] - 131s 23ms/step - loss: 0.6014 - accuracy: 0.8938 - val_loss: 0.6358 - val_accuracy: 0.8
Epoch 5/10
5626/5627 [=====>.] - ETA: 0s - loss: 0.5909 - accuracy: 0.8999
Epoch 5: ReduceLROnPlateau reducing learning rate to 0.00010000000474974513.
5627/5627 [=====] - 139s 25ms/step - loss: 0.5909 - accuracy: 0.8999 - val_loss: 0.6388 - val_accuracy: 0.8
Epoch 6/10
5627/5627 [=====] - 139s 25ms/step - loss: 0.5686 - accuracy: 0.9124 - val_loss: 0.6413 - val_accuracy: 0.8
Epoch 7/10

```



```
5626/5627 [=====>.] - ETA: 0s - loss: 0.5635 - accuracy: 0.9155
Epoch 7: ReduceLROnPlateau reducing learning rate to 1.0000000474974514e-05.
5627/5627 [=====] - 142s 25ms/step - loss: 0.5635 - accuracy: 0.9155 - val_loss: 0.6424 - val_accuracy: 0.8715
```

```
1 model.evaluate(test_ds)
```

```
942/942 [=====] - 16s 16ms/step - loss: 0.6438 - accuracy: 0.8715
[0.6438229084014893, 0.8715115189552307]
```

```
1 val_prob = model.predict(val_ds, verbose=1)
2 val_predictions = tf.argmax(val_prob, axis=1)
```

```
945/945 [=====] - 14s 14ms/step
```

```
1 test_prob = model.predict(test_ds, verbose=1)
2 test_predict = tf.argmax(test_prob, axis=1)
```

```
942/942 [=====] - 18s 19ms/step
```

```
1 print(val_predictions)
2 print(test_predict)
```

```
tf.Tensor([0 0 3 ... 4 1 1], shape=(30212,), dtype=int64)
tf.Tensor([0 2 2 ... 4 4 1], shape=(30135,), dtype=int64)
```

```
1 LE.classes_
```

```
array(['BACKGROUND', 'CONCLUSIONS', 'METHODS', 'OBJECTIVE', 'RESULTS'],
      dtype=object)
```

```
1 class_labels = LE.classes_
```

```
1 # Generate validation classification report
2 print(classification_report(val_labels_encoded, val_predictions, target_names=class_labels))
```

```
precision    recall  f1-score   support

BACKGROUND   0.71     0.82     0.76     3449
CONCLUSIONS 0.89     0.93     0.91     4582
METHODS       0.92     0.91     0.92     9964
OBJECTIVE     0.71     0.58     0.64     2376
RESULTS       0.92     0.90     0.91     9841

accuracy          0.87     30212
macro avg         0.83     0.83     0.83     30212
weighted avg      0.88     0.87     0.87     30212
```

```
1 # Generate test classification report
2 print(classification_report(test_labels_encoded, test_predict, target_names=class_labels))
```

```
precision    recall  f1-score   support

BACKGROUND   0.72     0.81     0.76     3621
CONCLUSIONS 0.89     0.93     0.91     4571
METHODS       0.91     0.92     0.92     9897
OBJECTIVE     0.69     0.59     0.63     2333
RESULTS       0.92     0.89     0.90     9713

accuracy          0.87     30135
macro avg         0.83     0.83     0.83     30135
weighted avg      0.87     0.87     0.87     30135
```

```

1 from sklearn.metrics import accuracy_score, top_k_accuracy_score, precision_recall_fscore_support, matthews_corrcoef
2 def performance_metrics(y_actual, y_predicted, y_probabilities):
3
4     model_accuracy = round(accuracy_score(y_actual, y_predicted), 5)
5     top_3_accuracy = round(top_k_accuracy_score(y_actual, y_probabilities, k=3), 5)
6     model_precision, model_recall, model_f1, _ = precision_recall_fscore_support(test_labels_encoded,
7                                                                               test_predict,
8                                                                               average="weighted")
9     model_precision, model_recall, model_f1 = round(model_precision, 5), round(model_recall, 5), round(model_f1, 5)
10    model_matthews_corrcoef = round(matthews_corrcoef(y_actual, y_predicted), 5)
11
12    print(f'\nPerformance Metrics:\n')
13    print('-----**-----')
14    print(f'accuracy_score:\t\t{model_accuracy}\n')
15    print('*****')
16    print(f'top_3_accuracy_score:\t\t{top_3_accuracy}\n')
17    print('*****')
18    print(f'precision_score:\t\t{model_precision}\n')
19    print('*****')
20    print(f'recall_score:\t\t{model_recall}\n')
21    print('*****')
22    print(f'f1_score:\t\t{model_f1}\n')
23    print('*****')
24    print(f'matthews_corrcoef:\t\t{model_matthews_corrcoef}\n')
25
26    return

```

```
1 performance_metrics(val_labels_encoded, val_predictions, val_prob)
```



Performance Metrics:

```

-----**-----
accuracy_score:          0.87452

*****
top_3_accuracy_score:    0.99533

*****
precision_score:         0.87183

*****
recall_score:            0.87151

*****
f1_score:                0.87094

*****
matthews_corrcoef:       0.83177

```

```
1 performance_metrics(test_labels_encoded, test_predict, test_prob)
```



Performance Metrics:

```

-----**-----
accuracy_score:          0.87151

*****
top_3_accuracy_score:    0.99363

*****
precision_score:         0.87183

*****
recall_score:            0.87151

*****
f1_score:                0.87094

*****
matthews_corrcoef:       0.82805

```


```
1 cm = confusion_matrix(test_labels_encoded, test_predict)
```

```

1 plt.figure(figsize=(10,10))
2 sns.heatmap(cm, annot=True, linewidths=1, xticklabels=class_labels, yticklabels=class_labels, fmt='d')

```

<Axes: >



```

1 # importing required libraries
2 import numpy as np
3 import pandas as pd
4 import os
5 import librosa
6 import librosa.display
7 import IPython as ipd
8 import glob
9
10 import tensorflow as tf
11 from tensorflow.keras.models import Sequential
12 from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten
13 from tensorflow.keras.layers import Convolution2D, Conv2D, MaxPooling2D, GlobalAveragePooling2D
14 from tensorflow.keras.optimizers import Adam
15 from sklearn import metrics
16 from tensorflow.keras.layers import Conv2D, Flatten, Dense, MaxPool2D, Dropout

```

```

1 # directory path of mounted gdrive
2 bas_dir_path = r'/content/drive/MyDrive/edge-collected-gunshot-audio.zip'

```

```

1 #creating directory for the extracted data
2 os.makedirs(os.path.join(os.getcwd(), 'extracted_data'))

```

```

1 # unzipping all the files into to extracted_data created directory
2 import zipfile
3 with zipfile.ZipFile(bas_dir_path, 'r') as zip_ref:
4     zip_ref.extractall(os.path.join(bas_dir_path, '/content/extracted_data'))

```

```
1 !pip install split-folders
```

Collecting split-folders
 Downloading split_folders-0.5.1-py3-none-any.whl (8.4 kB)
 Installing collected packages: split-folders
 Successfully installed split-folders-0.5.1

```

1 os.makedirs(os.path.join(os.getcwd(),'my_split_data')) # directory for storing the splitted data

1 # splitting the data into 80:20 ratio
2 import splitfolders
3 splitfolders.ratio('/content/extracted_data/edge-collected-gunshot-audio','/content/my_split_data',seed=42,ratio=(0.8,0.2))

↗ Copying files: 2148 files [00:11, 180.55 files/s]

1 train_dir_path = r'/content/my_split_data/train' # newly created train directory which contains 80% of the files
2 val_dir_path = r'/content/my_split_data/val' # newly created validation directory which contains 20% of the files

1 # creating function for extracting features of the given .wav files
2 def extract_features(filename):
3     audio, sr = librosa.load(filename,sr = None,res_type='kaiser_fast')
4     mfccs = librosa.feature.mfcc(y=audio,sr=sr,n_mfcc=40)
5     mfcc_scaled = np.mean(mfccs.T,axis=0)
6     return mfcc_scaled

1 # checking for sample_rate
2 ex_data, sampling_rate = librosa.load('/content/extracted_data/edge-collected-gunshot-audio/38s&ws_dot38_caliber/03fc4685-909e-42c5-')
3 print(sampling_rate)

↗ 44100

1 !pip install resampy # internally res_type='kaiser_fast' calling resampy

↗ Collecting resampy
  Downloading resampy-0.4.2-py3-none-any.whl (3.1 MB)
     3.1/3.1 MB 28.1 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from resampy) (1.23.5)
Requirement already satisfied: numba>=0.53 in /usr/local/lib/python3.10/dist-packages (from resampy) (0.58.1)
Requirement already satisfied: llvmlite<0.42,>=0.41.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba>=0.53->resampy) (0.42.0)
Installing collected packages: resampy
Successfully installed resampy-0.4.2

1 features=[] # collecting the features and labels for the .wav files
2 firearm = ''
3 for folder in os.listdir('/content/my_split_data/train'):
4     gen = glob.iglob('**/*.wav',root_dir=os.path.join('/content/my_split_data/train',folder),recursive=True) # using iglob generator for files in gen
5     for files in gen:
6         path = os.path.join(os.path.join('/content/my_split_data/train',folder),files) # path to provide for extract_features function.
7         data = extract_features(path) # calling extract_features functions for loading features
8         firearm = folder # assigning the name of the firearm as label
9         features.append([data,firearm]) # appending the extracted features and labels

1 #df = pd.DataFrame()
2 df = pd.DataFrame(features,columns=['feature','class']) # converting into dataframe for the extracted features and labels
3 df.head()

↗

```

	feature	class
0	[-479.21558, 115.43387, 4.464393, -2.232155, 0...	remington_870_12_gauge
1	[-622.7922, 160.48073, -31.198713, -13.155003,...	remington_870_12_gauge
2	[-685.34485, 205.09016, -12.939822, -16.207764...	remington_870_12_gauge
3	[-244.56116, 180.25233, -31.406918, 10.96077, ...	remington_870_12_gauge
4	[-683.1073, 138.53195, 10.358094, -3.0826027, ...	remington_870_12_gauge

```

1 X_copy = df['feature'].copy() # copying the data from the above dataframe for not editing the original data.
2 y_copy = df['class'].copy() # copying the labels correspondingly.

1 X = np.array(X_copy.to_list()) # into list
2 y = np.array(y_copy.to_list())

1 X = np.array(X_copy.to_list()) # converting into array
2 y = np.array(y_copy.to_list())
3 from sklearn.preprocessing import LabelEncoder
4 from keras.utils import to_categorical
5 LE = LabelEncoder()
6 yy = to_categorical(LE.fit_transform(y)) # fitting and transforming on training data of labels and converting into to_categorical
7 from sklearn.model_selection import train_test_split
8 X_train,X_test,y_train,y_test = train_test_split(X,yy, test_size=0.2, random_state=42) # splitting into train and test

```

```

1 #validation data
2 val_features=[] # applying fucntion on validation data and storing in list.
3 val_labels = ''
4 for folder in os.listdir('/content/my_split_data/val'):
5     gen = glob.iglob('**/*.wav',root_dir=os.path.join('/content/my_split_data/val',folder),recursive=True)
6     for files in gen:
7         path = os.path.join(os.path.join('/content/my_split_data/val',folder),files)
8         value = extract_features(path)
9         val_labels = folder
10        val_features.append([value,val_labels])

1 # total number of files in validation data
2 len(os.listdir('/content/my_split_data/val/38s&ws_dot38_caliber')) + len(os.listdir('/content/my_split_data/val/glock_17_9mm_caliber

431

1 val_df = pd.DataFrame(val_features,columns=['val_features','val_labels']) # into DataFrame

1 val_df.shape # verifying the shape of the data

(431, 2)

1 x_val_copy = val_df['val_features'].copy() # creating copies of original data
2 y_val_copy = val_df['val_labels'].copy()
3 x_val_array = np.array(x_val_copy.to_list()) # into list
4 y_val_array = np.array(y_val_copy.to_list())

1 y_val_array_encoded = to_categorical(LE.fit_transform(y_val_array)) # transforming with the help of 'LE' and convering into to_categ

1 num_rows = 4
2 num_columns = 10
3 num_channels = 1
4 num_labels = yy.shape[1]
5 filter_size = 2
6 x_train = X_train.reshape(X_train.shape[0], num_rows, num_columns, num_channels) # reshaping the training and validation data for fe
7 x_val = x_val_array.reshape(x_val_array.shape[0], num_rows, num_columns, num_channels)

1 # Construct model
2 model = Sequential()
3 model.add(Conv2D(filters=16, kernel_size=2, padding="same", input_shape=(num_rows, num_columns, num_channels), activation='relu'))
4 model.add(MaxPooling2D(pool_size=1))
5 model.add(Dropout(0.2))
6
7 model.add(Conv2D(filters=32, kernel_size=2, padding="same", activation='relu'))
8 model.add(MaxPooling2D(pool_size=1))
9 model.add(Dropout(0.2))
10
11 model.add(Conv2D(filters=64, kernel_size=2, padding="same", activation='relu'))
12 model.add(MaxPooling2D(pool_size=1))
13 model.add(Dropout(0.2))
14
15 model.add(Conv2D(filters=128, kernel_size=2, padding="same", activation='relu'))
16 model.add(MaxPooling2D(pool_size=1))
17 model.add(Dropout(0.2))
18 model.add(GlobalAveragePooling2D())
19
20 model.add(Dense(num_labels, activation='softmax'))

1 # Compile the model
2 model.compile(loss='binary_crossentropy', metrics=['accuracy'], optimizer='adam')
3
4 # Display model architecture summary
5 model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 4, 10, 16)	80
max_pooling2d_4 (MaxPoolin g2D)	(None, 4, 10, 16)	0
dropout_4 (Dropout)	(None, 4, 10, 16)	0
conv2d_5 (Conv2D)	(None, 4, 10, 32)	2080

max_pooling2d_5 (MaxPoolin g2D)	(None, 4, 10, 32)	0
dropout_5 (Dropout)	(None, 4, 10, 32)	0
conv2d_6 (Conv2D)	(None, 4, 10, 64)	8256
max_pooling2d_6 (MaxPoolin g2D)	(None, 4, 10, 64)	0
dropout_6 (Dropout)	(None, 4, 10, 64)	0
conv2d_7 (Conv2D)	(None, 4, 10, 128)	32896
max_pooling2d_7 (MaxPoolin g2D)	(None, 4, 10, 128)	0
dropout_7 (Dropout)	(None, 4, 10, 128)	0
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 128)	0
dense_1 (Dense)	(None, 4)	516

```

=====
Total params: 43828 (171.20 KB)
Trainable params: 43828 (171.20 KB)
Non-trainable params: 0 (0.00 Byte)

```

```

1 num_epochs = 50
2 num_batch_size = 50
3 # fitting the model
4 model.fit(x_train, y_train, batch_size=num_batch_size, epochs=num_epochs, validation_data=(x_val, y_val_array_encoded), verbose=1)

```

```

Epoch 1/50
28/28 [=====] - 4s 68ms/step - loss: 0.5877 - accuracy: 0.2935 - val_loss: 0.5705 - val_accuracy: 0.310
Epoch 2/50
28/28 [=====] - 2s 54ms/step - loss: 0.5601 - accuracy: 0.3081 - val_loss: 0.5588 - val_accuracy: 0.324
Epoch 3/50
28/28 [=====] - 1s 35ms/step - loss: 0.5561 - accuracy: 0.3248 - val_loss: 0.5538 - val_accuracy: 0.368
Epoch 4/50
28/28 [=====] - 1s 34ms/step - loss: 0.5531 - accuracy: 0.3372 - val_loss: 0.5491 - val_accuracy: 0.350
Epoch 5/50
28/28 [=====] - 1s 33ms/step - loss: 0.5478 - accuracy: 0.3540 - val_loss: 0.5449 - val_accuracy: 0.385
Epoch 6/50
28/28 [=====] - 1s 33ms/step - loss: 0.5407 - accuracy: 0.4020 - val_loss: 0.5401 - val_accuracy: 0.452
Epoch 7/50
28/28 [=====] - 1s 33ms/step - loss: 0.5420 - accuracy: 0.3889 - val_loss: 0.5352 - val_accuracy: 0.431
Epoch 8/50
28/28 [=====] - 1s 33ms/step - loss: 0.5355 - accuracy: 0.4086 - val_loss: 0.5297 - val_accuracy: 0.457
Epoch 9/50
28/28 [=====] - 1s 33ms/step - loss: 0.5259 - accuracy: 0.4428 - val_loss: 0.5193 - val_accuracy: 0.457
Epoch 10/50
28/28 [=====] - 1s 36ms/step - loss: 0.5183 - accuracy: 0.4363 - val_loss: 0.5173 - val_accuracy: 0.422
Epoch 11/50
28/28 [=====] - 1s 34ms/step - loss: 0.5165 - accuracy: 0.4319 - val_loss: 0.5087 - val_accuracy: 0.443
Epoch 12/50
28/28 [=====] - 1s 43ms/step - loss: 0.5121 - accuracy: 0.4443 - val_loss: 0.5049 - val_accuracy: 0.475
Epoch 13/50
28/28 [=====] - 1s 44ms/step - loss: 0.5094 - accuracy: 0.4545 - val_loss: 0.4965 - val_accuracy: 0.489
Epoch 14/50
28/28 [=====] - 2s 55ms/step - loss: 0.5003 - accuracy: 0.4705 - val_loss: 0.4967 - val_accuracy: 0.466
Epoch 15/50
28/28 [=====] - 1s 51ms/step - loss: 0.4991 - accuracy: 0.4727 - val_loss: 0.4868 - val_accuracy: 0.508
Epoch 16/50
28/28 [=====] - 2s 54ms/step - loss: 0.4996 - accuracy: 0.4639 - val_loss: 0.5043 - val_accuracy: 0.433
Epoch 17/50
28/28 [=====] - 2s 54ms/step - loss: 0.4952 - accuracy: 0.4843 - val_loss: 0.4940 - val_accuracy: 0.482
Epoch 18/50
28/28 [=====] - 1s 46ms/step - loss: 0.4932 - accuracy: 0.4727 - val_loss: 0.4835 - val_accuracy: 0.487
Epoch 19/50
28/28 [=====] - 1s 37ms/step - loss: 0.4879 - accuracy: 0.4953 - val_loss: 0.4790 - val_accuracy: 0.508
Epoch 20/50
28/28 [=====] - 1s 33ms/step - loss: 0.4835 - accuracy: 0.4982 - val_loss: 0.4718 - val_accuracy: 0.524
Epoch 21/50
28/28 [=====] - 1s 33ms/step - loss: 0.4885 - accuracy: 0.5011 - val_loss: 0.4758 - val_accuracy: 0.512
Epoch 22/50
28/28 [=====] - 1s 33ms/step - loss: 0.4764 - accuracy: 0.5157 - val_loss: 0.4851 - val_accuracy: 0.489
Epoch 23/50
28/28 [=====] - 1s 33ms/step - loss: 0.4775 - accuracy: 0.5025 - val_loss: 0.4679 - val_accuracy: 0.515
Epoch 24/50
28/28 [=====] - 1s 33ms/step - loss: 0.4703 - accuracy: 0.5178 - val_loss: 0.4595 - val_accuracy: 0.549
Epoch 25/50
28/28 [=====] - 1s 32ms/step - loss: 0.4704 - accuracy: 0.5310 - val_loss: 0.4723 - val_accuracy: 0.536
Epoch 26/50
28/28 [=====] - 1s 35ms/step - loss: 0.4677 - accuracy: 0.5302 - val_loss: 0.4570 - val_accuracy: 0.584
Epoch 27/50
28/28 [=====] - 1s 34ms/step - loss: 0.4648 - accuracy: 0.5302 - val_loss: 0.4628 - val_accuracy: 0.568

```

```

Epoch 28/50
28/28 [=====] - 1s 32ms/step - loss: 0.4615 - accuracy: 0.5310 - val_loss: 0.4684 - val_accuracy: 0.538
Epoch 29/50

1 X_test = X_test.reshape(X_test.shape[0], num_rows, num_columns, num_channels) # converting test_data for predictions

1 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
2 from tensorflow.math import confusion_matrix
3 # Evaluate the model
4 y_pred = model.predict(X_test)
5 y_pred_classes = np.argmax(y_pred, axis=1)
6 y_test_classes = np.argmax(y_test, axis=1)
7
8 # Calculate evaluation metrics
9 accuracy = accuracy_score(y_test_classes, y_pred_classes)
10 precision = precision_score(y_test_classes, y_pred_classes, average='weighted')
11 recall = recall_score(y_test_classes, y_pred_classes, average='weighted')
12 f1 = f1_score(y_test_classes, y_pred_classes, average='weighted')
13
14 # Print the evaluation metrics
15 print(f'Accuracy - CNN: {accuracy}')
16 print(f'Precision - CNN: {precision}')
17 print(f'Recall - CNN: {recall}')
18 print(f'F1 Score - CNN: {f1}')
19
20 print('Confusion_matrix - CNN: ',confusion_matrix(y_test_classes, y_pred_classes))

11/11 [=====] - 0s 22ms/step
Accuracy - CNN: 0.6133720930232558
Precision - CNN: 0.6331407042803696
Recall - CNN: 0.6133720930232558
F1 Score - CNN: 0.5888214198046698
Confusion_matrix - CNN: tf.Tensor(
[[35 12  3 22]
 [14 73  1 24]
 [10 18 14 19]
 [ 2  7  1 89]], shape=(4, 4), dtype=int32)

1 from tensorflow.keras.layers import LSTM, Dense,Dropout,Flatten
2 from tensorflow import keras
3 # Build a simple LSTM model
4 model_LSTM = Sequential()
5 model_LSTM.add(LSTM(128, input_shape=(X_train.shape[1], 1)))
6 model_LSTM.add(Dropout(0.5))
7 model_LSTM.add(Dense(len(LE.classes_), activation='softmax'))
8
9 # Compile the model
10 model_LSTM.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
11
12 # Reshape features for LSTM input
13 x_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))
14 x_val = x_val_array.reshape((x_val_array.shape[0], x_val_array.shape[1], 1))
15 x_test = X_test.reshape((X_test.shape[0], X_test.shape[1], 1))
16
17 # Train the model
18 model_LSTM.fit(x_train, y_train, epochs=50, batch_size=32, validation_data=(x_val, y_val_array_encoded), verbose=1)
19
20 # Evaluate the model
21 loss, accuracy = model_LSTM.evaluate(x_test, y_test)
22 print(f'Test Loss: {loss}, Test Accuracy: {accuracy}')
23

Epoch 1/50
43/43 [=====] - 8s 74ms/step - loss: 1.3604 - accuracy: 0.3132 - val_loss: 1.3335 - val_accuracy: 0.348
Epoch 2/50
43/43 [=====] - 4s 85ms/step - loss: 1.3104 - accuracy: 0.3846 - val_loss: 1.2868 - val_accuracy: 0.375
Epoch 3/50
43/43 [=====] - 4s 102ms/step - loss: 1.2635 - accuracy: 0.4246 - val_loss: 1.2519 - val_accuracy: 0.43
Epoch 4/50
43/43 [=====] - 3s 62ms/step - loss: 1.2567 - accuracy: 0.4275 - val_loss: 1.2413 - val_accuracy: 0.436
Epoch 5/50
43/43 [=====] - 2s 56ms/step - loss: 1.2228 - accuracy: 0.4545 - val_loss: 1.2253 - val_accuracy: 0.431
Epoch 6/50
43/43 [=====] - 2s 55ms/step - loss: 1.2218 - accuracy: 0.4712 - val_loss: 1.2191 - val_accuracy: 0.468
Epoch 7/50
43/43 [=====] - 4s 97ms/step - loss: 1.1966 - accuracy: 0.4669 - val_loss: 1.1939 - val_accuracy: 0.466
Epoch 8/50
43/43 [=====] - 4s 103ms/step - loss: 1.1774 - accuracy: 0.4894 - val_loss: 1.2091 - val_accuracy: 0.43
Epoch 9/50
43/43 [=====] - 3s 76ms/step - loss: 1.1400 - accuracy: 0.5040 - val_loss: 1.1334 - val_accuracy: 0.505
Epoch 10/50
43/43 [=====] - 2s 56ms/step - loss: 1.1336 - accuracy: 0.5127 - val_loss: 1.1332 - val_accuracy: 0.522
Epoch 11/50

```

```

43/43 [=====] - 2s 57ms/step - loss: 1.1263 - accuracy: 0.5222 - val_loss: 1.1022 - val_accuracy: 0.522
Epoch 12/50
43/43 [=====] - 2s 56ms/step - loss: 1.0958 - accuracy: 0.5317 - val_loss: 1.0851 - val_accuracy: 0.561
Epoch 13/50
43/43 [=====] - 3s 76ms/step - loss: 1.0777 - accuracy: 0.5317 - val_loss: 1.0578 - val_accuracy: 0.540
Epoch 14/50
43/43 [=====] - 4s 96ms/step - loss: 1.0772 - accuracy: 0.5433 - val_loss: 1.0943 - val_accuracy: 0.517
Epoch 15/50
43/43 [=====] - 3s 71ms/step - loss: 1.0912 - accuracy: 0.5193 - val_loss: 1.0740 - val_accuracy: 0.540
Epoch 16/50
43/43 [=====] - 2s 54ms/step - loss: 1.0460 - accuracy: 0.5564 - val_loss: 1.0626 - val_accuracy: 0.554
Epoch 17/50
43/43 [=====] - 2s 54ms/step - loss: 1.0217 - accuracy: 0.5717 - val_loss: 1.0229 - val_accuracy: 0.554
Epoch 18/50
43/43 [=====] - 2s 55ms/step - loss: 0.9757 - accuracy: 0.5870 - val_loss: 0.9769 - val_accuracy: 0.589
Epoch 19/50
43/43 [=====] - 3s 68ms/step - loss: 0.9680 - accuracy: 0.5972 - val_loss: 1.0552 - val_accuracy: 0.536
Epoch 20/50
43/43 [=====] - 4s 102ms/step - loss: 0.9286 - accuracy: 0.6111 - val_loss: 0.9865 - val_accuracy: 0.58
Epoch 21/50
43/43 [=====] - 5s 112ms/step - loss: 0.9033 - accuracy: 0.6242 - val_loss: 0.9597 - val_accuracy: 0.60
Epoch 22/50
43/43 [=====] - 2s 55ms/step - loss: 0.9037 - accuracy: 0.6373 - val_loss: 0.9576 - val_accuracy: 0.596
Epoch 23/50
43/43 [=====] - 2s 53ms/step - loss: 0.8743 - accuracy: 0.6431 - val_loss: 0.9737 - val_accuracy: 0.589
Epoch 24/50
43/43 [=====] - 2s 55ms/step - loss: 0.8083 - accuracy: 0.6686 - val_loss: 0.8655 - val_accuracy: 0.619
Epoch 25/50
43/43 [=====] - 4s 96ms/step - loss: 0.8155 - accuracy: 0.6715 - val_loss: 0.9576 - val_accuracy: 0.573
Epoch 26/50
43/43 [=====] - 4s 101ms/step - loss: 0.7919 - accuracy: 0.6934 - val_loss: 0.9172 - val_accuracy: 0.62
Epoch 27/50
43/43 [=====] - 2s 57ms/step - loss: 0.7621 - accuracy: 0.7087 - val_loss: 0.9094 - val_accuracy: 0.638
Epoch 28/50
43/43 [=====] - 2s 55ms/step - loss: 0.7228 - accuracy: 0.7109 - val_loss: 0.8931 - val_accuracy: 0.628
Epoch 29/50

```

```

1 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
2
3 # Evaluate the model
4 y_pred = model_LSTM.predict(x_test)
5 y_pred_classes = np.argmax(y_pred, axis=1)
6 y_test_classes = np.argmax(y_test, axis=1)
7
8 # Calculate evaluation metrics
9 accuracy = accuracy_score(y_test_classes, y_pred_classes)
10 precision = precision_score(y_test_classes, y_pred_classes, average='weighted')
11 recall = recall_score(y_test_classes, y_pred_classes, average='weighted')
12 f1 = f1_score(y_test_classes, y_pred_classes, average='weighted')
13
14 # Print the evaluation metrics
15 print(f'Accuracy - LSTM: {accuracy}')
16 print(f'Precision - LSTM: {precision}')

```